



PPar Lunch 2014/12/10

Iterative Compilation 1.01

Compiler optimisations

- Modern Compilers:
 - 100s of possible code transformations
- GCC 4.8.1:
 - 205 optimisation options
 - 138 parameters affecting optimisation

Compiler optimisations

— $\sim 10^{340}$ combinations!

— $\sim 10^{82}$ atoms in the universe




Atoms in
Observable
Universe

GCC
Optimisation
Space

Compiler optimisations

— $\sim 10^{340}$ combinations!

— $\sim 10^{82}$ atoms in the universe






How do we choose optimisations?

How do we choose optimisations?

- Try each one on its own
 - Transformations affect each other
- Use an expert to find the best combo
 - Experts are few
 - One optimisation set doesn't fit all
- Use analytical models
 - Could help, but far from perfect
 - Architecture dependent


Choosing options - Example

- Usual approach: Let's pretend -O2 is the optimal choice
 - Enables transformations which usually improve performance
 - But not the optimal set of transformations for each program
- libquantum benchmark:



Choosing options - Example

- Usual approach: Let's pretend -O2 is the optimal choice
 - Enables transformations which usually improve performance
 - But not the optimal set of transformations for each program
- libquantum benchmark:

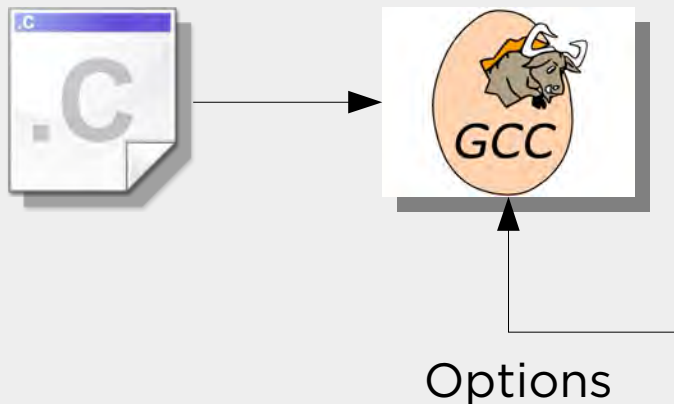


Choosing options

- Performance is lost because of inefficient optimisation strategies
- How do we harvest all the potential performance?

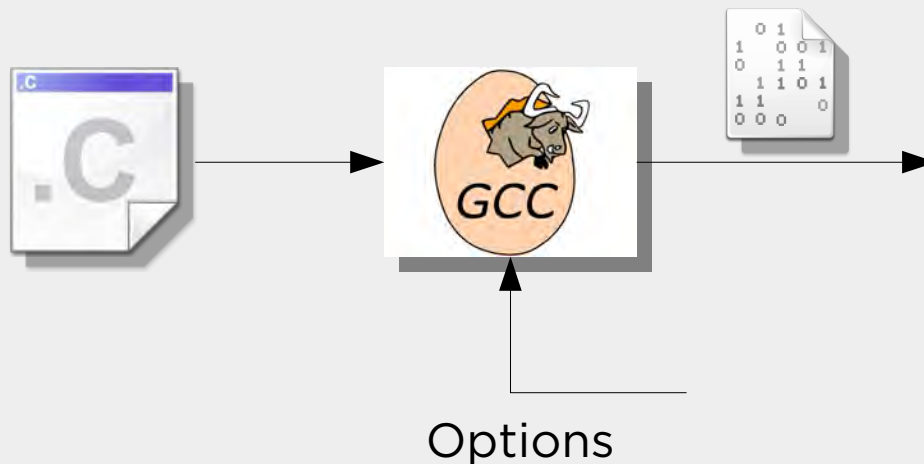
Iterative compilation

- Choose options



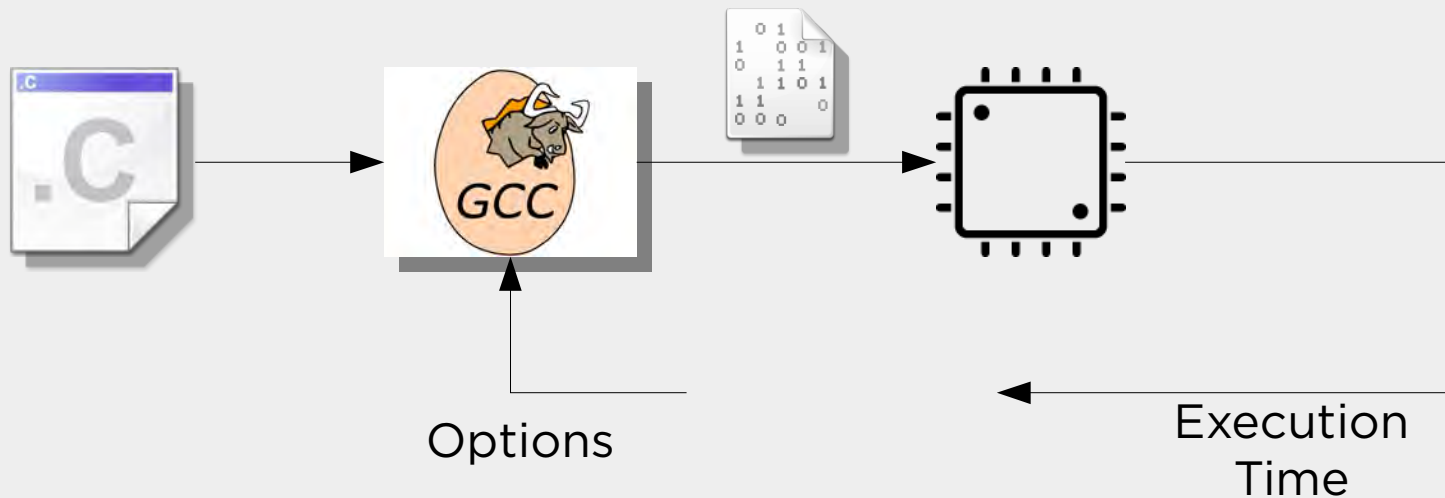
Iterative compilation

- Choose options
- Build the executables



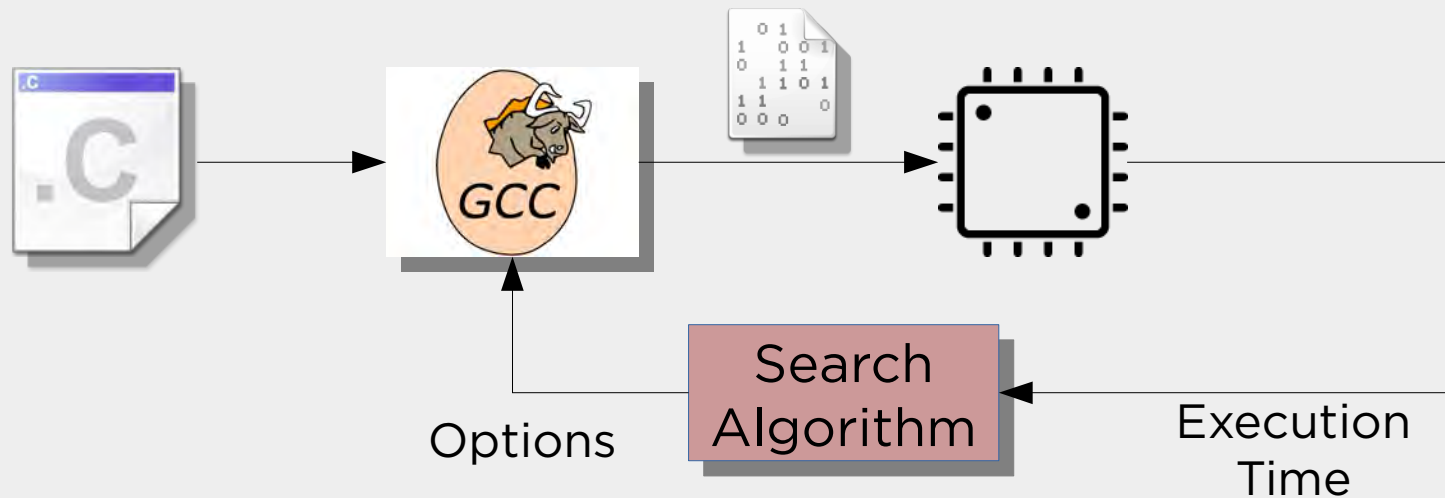
Iterative compilation

- Choose options
- Build the executables
- Get their runtime



Iterative compilation

- Choose options
- Build the executables
- Get their runtime
- Based on the results, choose the next round of optimisation options and repeat



Search Algorithms


- Search algorithm choice → critical parameter
 - Genetic Algorithm, Simulated Annealing, Grid Search, Window Search, etc.
- General Idea:
 - Find high performing areas of the space
 - Improve your knowledge of these areas




Iterative compilation - Example

- libquantum + Genetic Algorithm


Iterative compilation - Example




Iterative compilation - Example




Iterative compilation - Example




Iterative compilation - Example



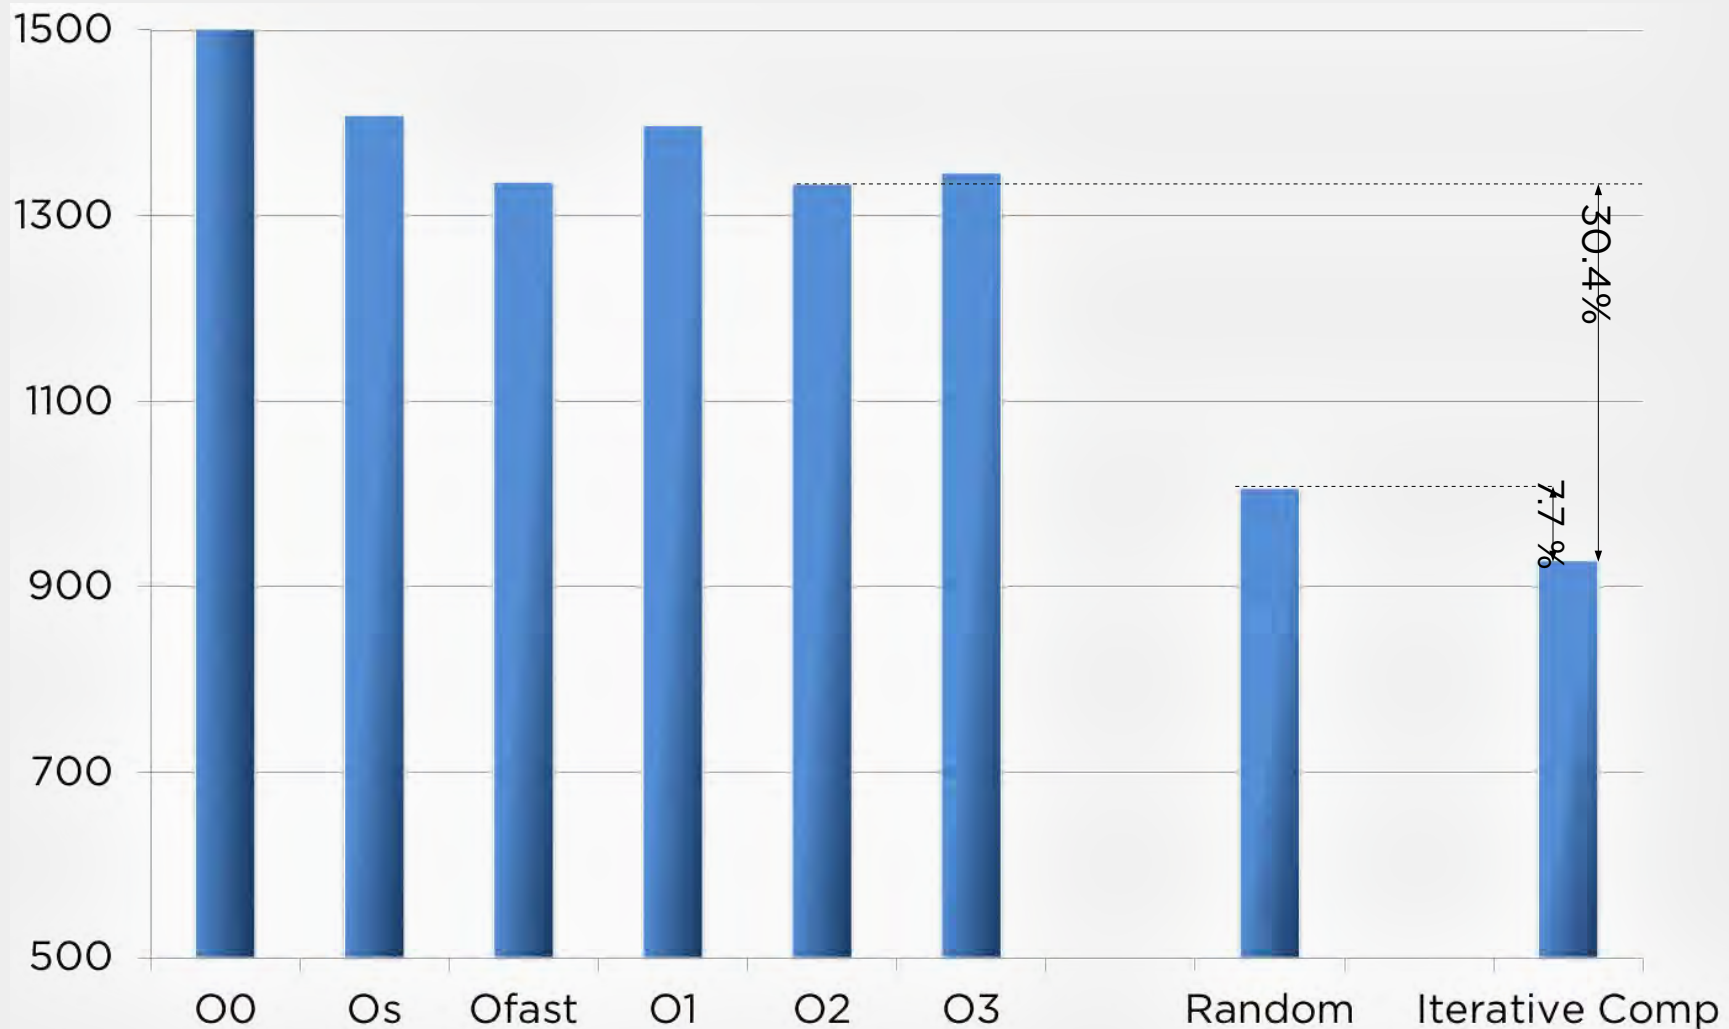
Iterative compilation - Example



Iterative compilation - Example



Iterative compilation - Example



Caveats


- Local maxima
- Noise
- Compiler errors

Local Maxima

- Wide & shallow vs narrow & deep search
 - Wide and shallow search not bothered by local maxima, but slow
 - Narrow and deep search focuses on local maxima, but might miss the optimal
- Algorithms and their parameters chosen to control this trade-off

Local Maxima - Example

- GA with tournament selection
 - Various tournament selection probabilities



Noise

- Noise makes iterative compilation less efficient


Noise

- Noise makes iterative compilation less efficient
- If noise comparable to the performance variation in the population
 - search algorithm becomes random


Noise - Example

- libquantum + gaussian noise

Pop. Average



Pop. Best


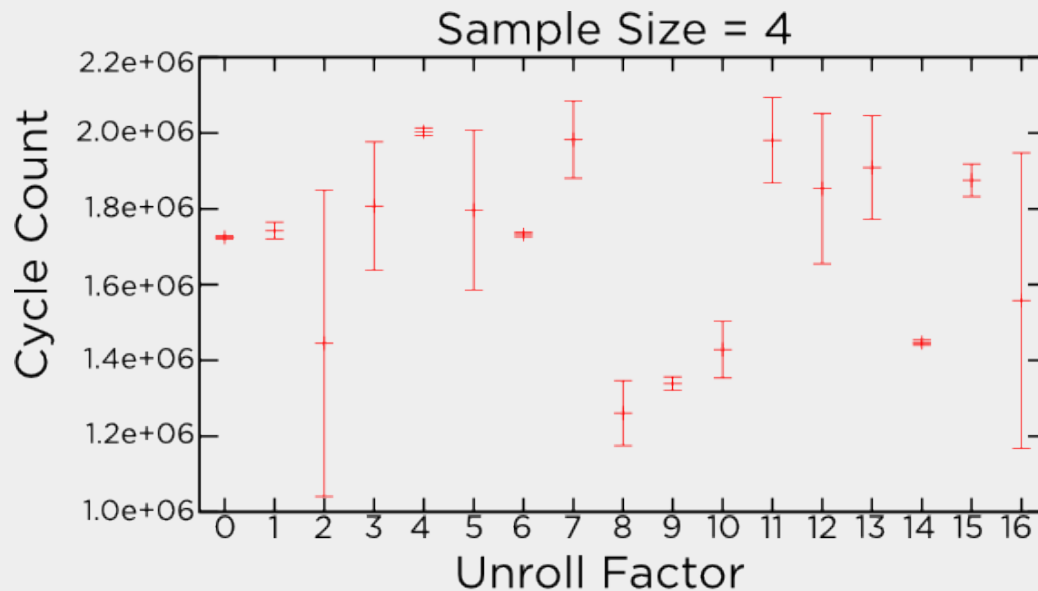


Noise

- Multiple observations to establish with some certainty the runtime of an executable
 - Higher evaluation overhead
 - slow

Noise



- Do we need high certainty?



Noise

- Do we need high certainty?


Do we care?



Noise

- Do we need high certainty?

Not really



Compiler Errors

- By randomly combining optimisations we enter an uncharted area
- Some combinations crash the compiler:
 - They don't work together and we know it
 - They don't work together and we don't know it
 - They trigger unknown bugs in the compiler code
- Crashes affect the speed of iterative compilation
 - Time wasted on unsuccessful builds
 - Benefit if we find problematic combinations early

Compiler Errors

- Others affect the program correctness
 - Due to bugs or unsafe transformations
- How do we handle them?
 - Fully deterministic behaviour
 - Output sensitive on changes in the control flow
 - Validation of the output after each execution

Improving iterative compilation

- Focus searching on areas predicted to be near optimal
 - Human expert input
 - Analytical models
- Pruning of irrelevant/harmful transformations
- Reduce evaluation overhead
 - Iterative compilation on the function call level



Iterative compilation - Drawbacks

Iterative compilation - Drawbacks

- Takes too long
 - Hours to month for the search to terminate
- Search parameters might need to be hand tuned
- Still, it's worth the effort for:
 - embedded systems
 - libraries
 - and others

Beyond iterative compilation

- Most of the cost → evaluation
 - 1000s of evaluations
 - To the garbage bin when we complete the search
- Can we somehow reuse the evaluation results for different programs?

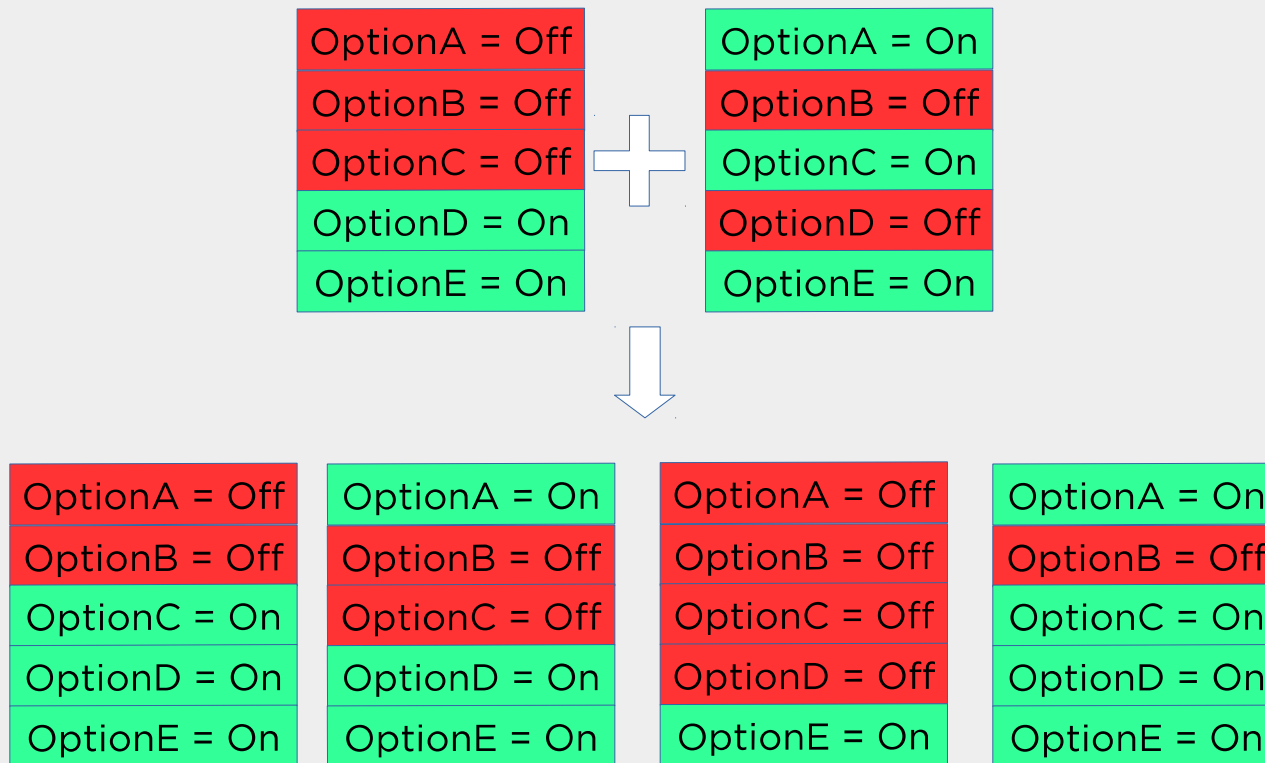


THE UNIVERSITY *of* EDINBURGH
informatics

icsa

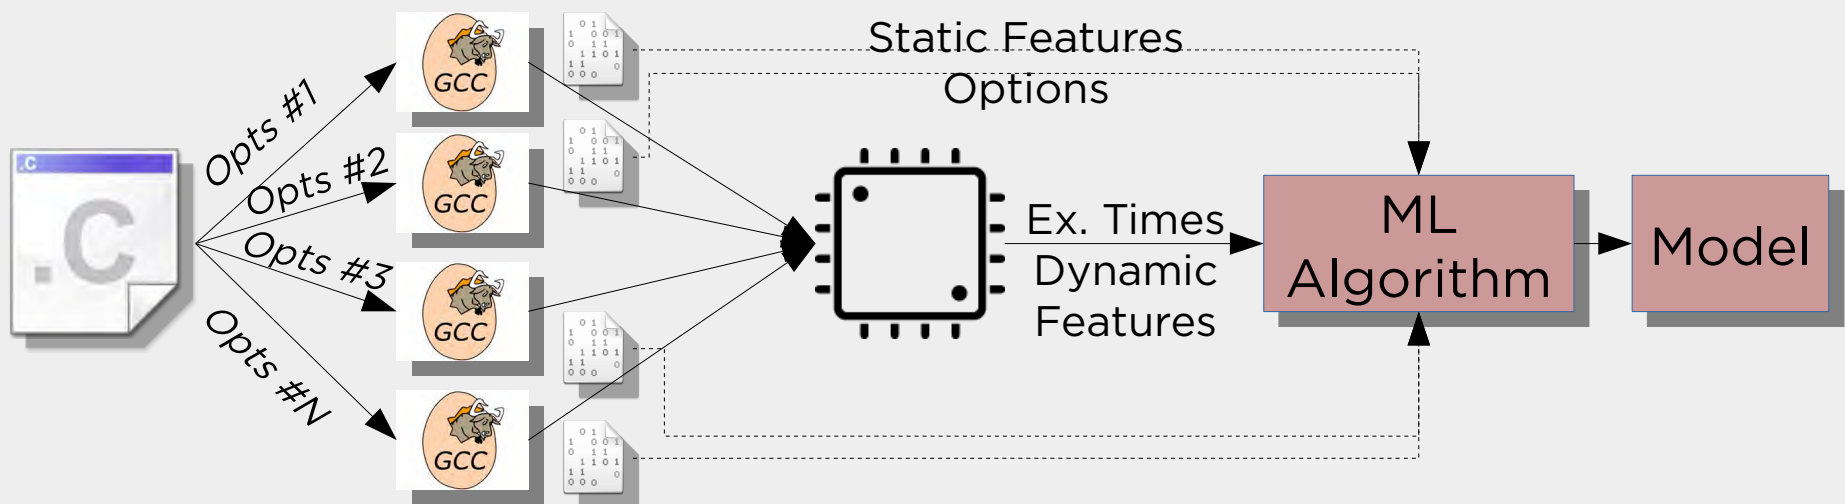
Institute for Computing
Systems Architecture

Genetic Algorithm - Example




ML-based optimising compiler

- Choose options
- Build
- Evaluate
- Train



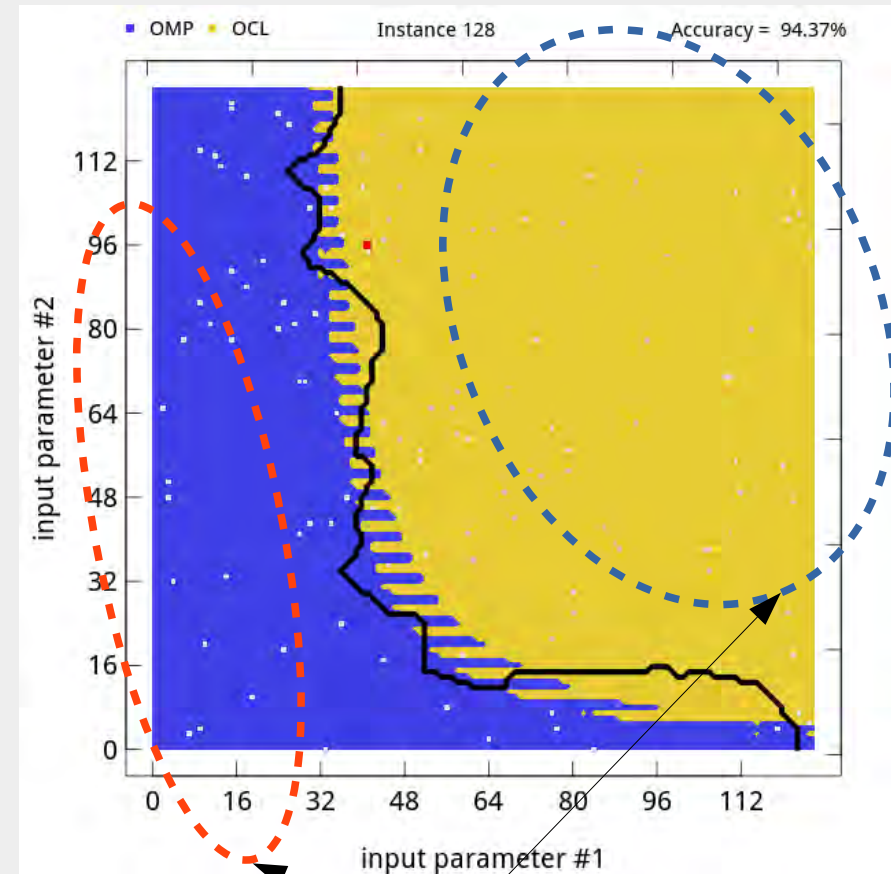
Example

- Best device (cpu or gpu) for an app based on the input size
 - Hotspot benchmark, 2d input
 - Random input sizes evaluated
 - Model built with the nearest neighbour algorithm



Example

- Most points offer little to no information → don't improve the model




Waste of time

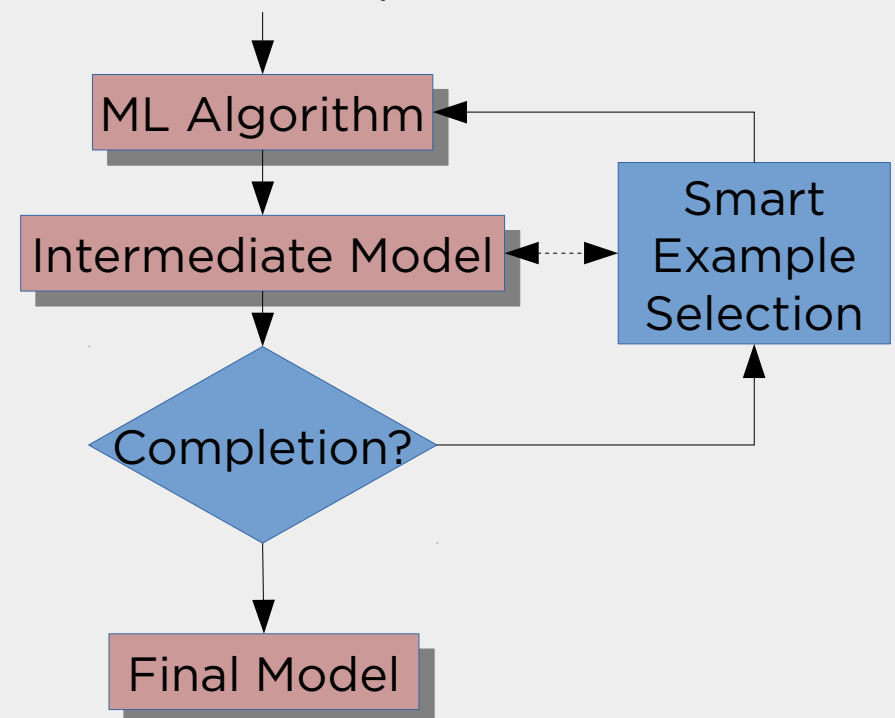
Active Learning

- Random training points selection wasteful
- Active Learning

Thousand of
random examples




few
random examples

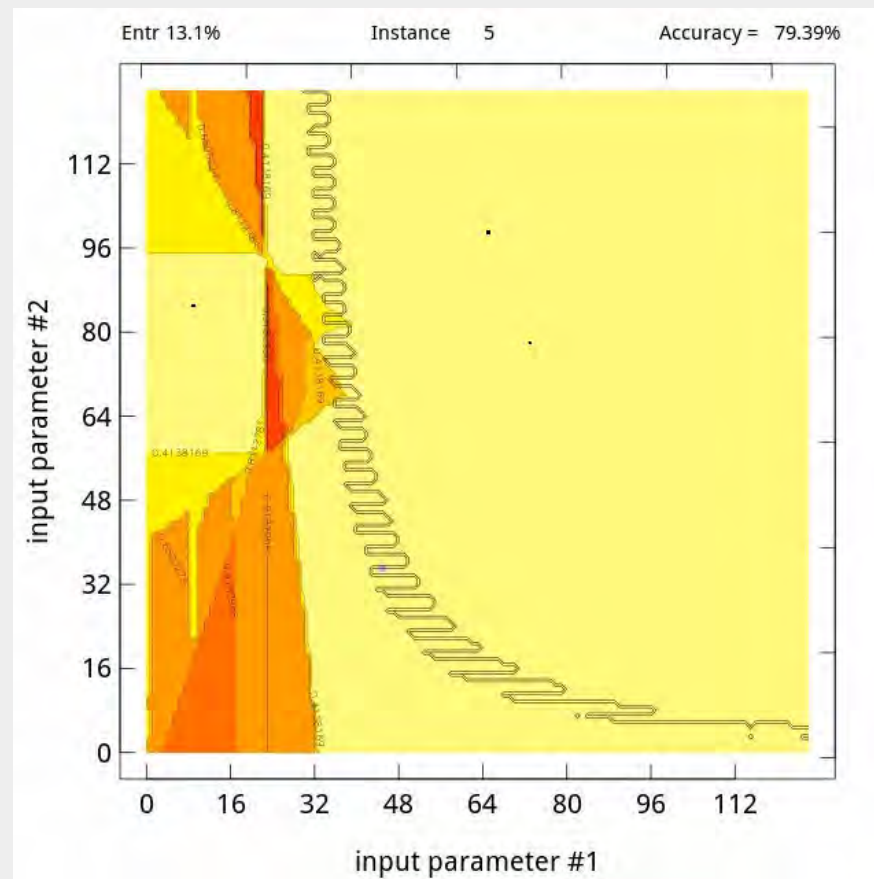


Active Learning

- Best points → High information points
 - close to the cpu-gpu border
 - max uncertainty about the best device
- Quantifying uncertainty:
 - Multiple ML algorithms trained with the same points
 - Uncertainty == degree of disagreement



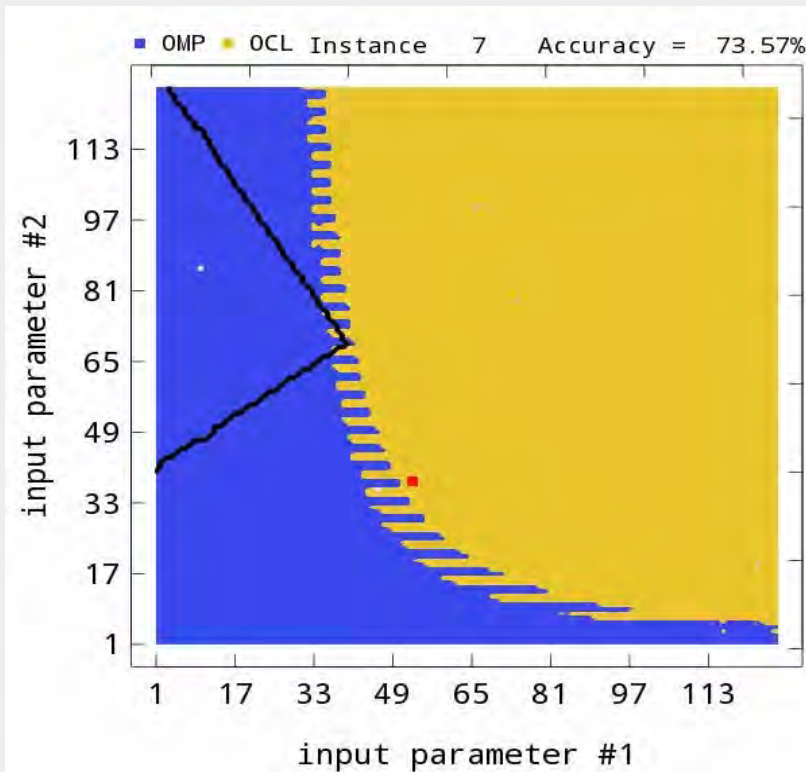
Disagreement



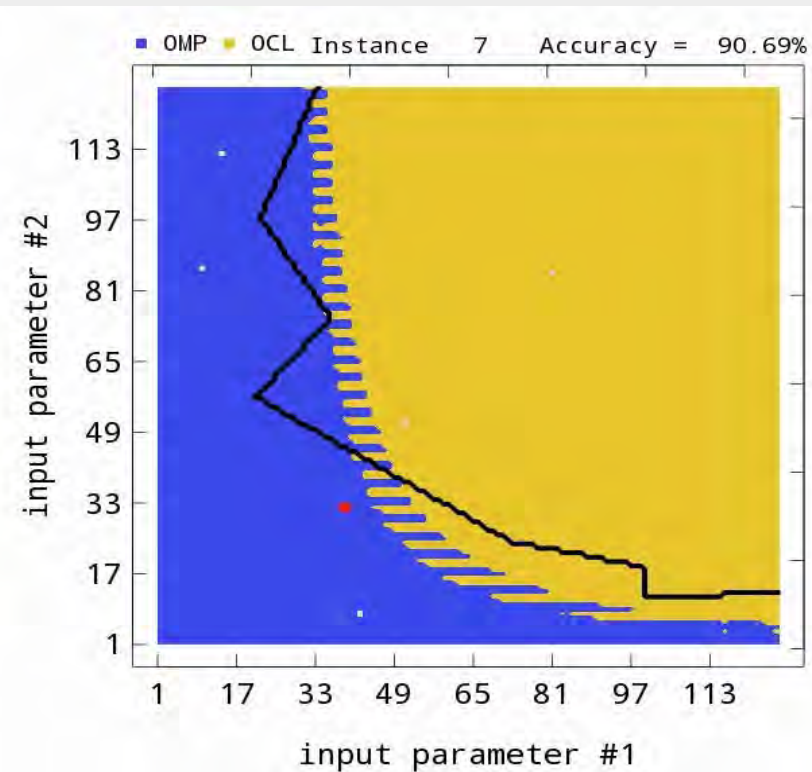
Comparison

Random

Active Learning



130 training points
for 95% accuracy



85 training points
for 95% accuracy



THE UNIVERSITY *of* EDINBURGH
informatics

icsa

Institute for Computing
Systems Architecture

Iterative compilation

- Choose options
- Build the executables
- Get their runtime
- Based on the results, choose the next round of optimisation options and repeat

